


Blockkurs:
Introduction to Machine Learning for Psychologists

Introduction to Machine Learning (and R)



Yannick Rothacher

Zürich, FS2025

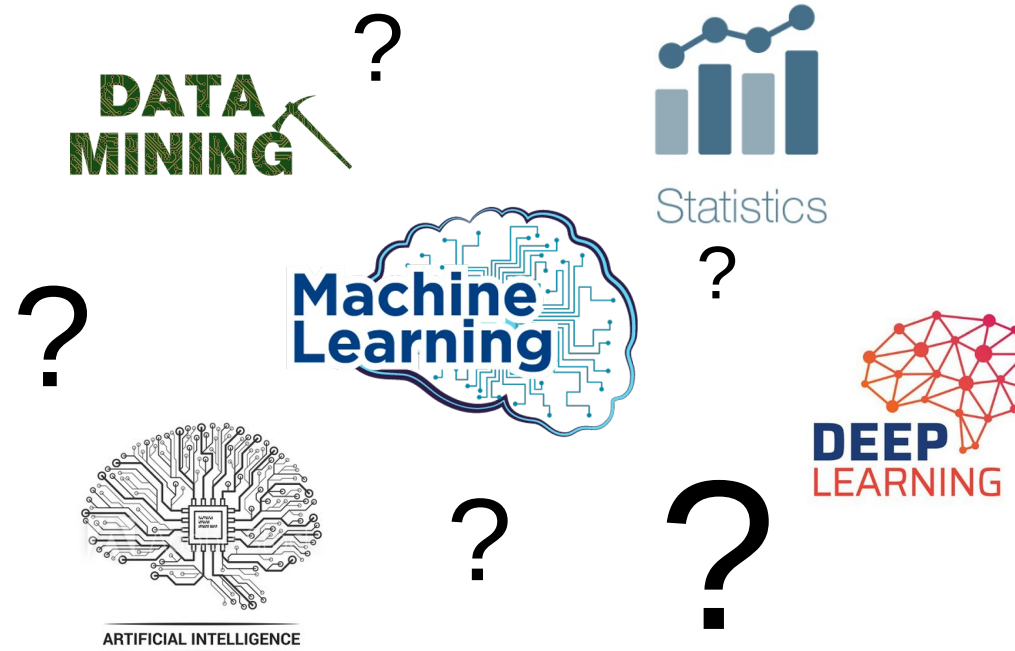
Who am I?

- ▶ **Yannick Rothacher**
- ▶ Background: Biology
(PhD in Neuroscience)
- ▶ Further education in "applied statistics" at ETH
Zürich
- ▶ Post-Doc at the Professorship for Psychological
Methods, Evaluation and Statistics (Prof. Carolin
Strobl)
- ▶ Hired as "Data Scientist" at Swiss Paraplegic
Research
- ▶ yrothacher@gmail.com



Course organization

- ▶ Two day course
- ▶ Mixture of lectures and practical exercises in R
- ▶ To get the credit point you have to write an **analysis report** after the course
 - ▶ The idea is that you take a data set ideally from your research, and apply one or multiple methods from this course to it
 - ▶ If you do not have access to a suitable data set I can provide you with one
 - ▶ The report has to be **handed in per mail until the 27. April 2025**
- ▶ Material is available on:
<https://github.com/ryannick28/MLCourse2025>




In groups:

- ▶ Why are you interested in machine learning?
- ▶ What are your expectations of this course?
- ▶ What do you associate with the term “machine learning”?

Course goals

- ▶ Give an insight into **various methods** in Machine Learning
- ▶ Teach the operating principles of the presented algorithms
- ▶ Practice the application of Machine Learning methods to data
- ▶ Deepen your skills in **R**



Tentative timetable

--- DAY 1 (13.3, AND-3-46) ---

09:30 - 11:00 Welcome + RIntro

11:00 - 11:45 PCA

11:45 - 12:30 PCA Exercise

--Lunch--

14:00 - 14:45 K-Means

14:45 - 15:15 K-Means Exercise

14:15 - 15:45 KNN

15:45 - 16:15 KNN Exercise

16:15 - 17:15 Crossvalidation + Write own function

--- DAY 2 (14.3, AFL-E-009) ---

09:30 - 10:30 Decision trees

10:30 - 11:30 Decision trees Exercise

--Lunch--

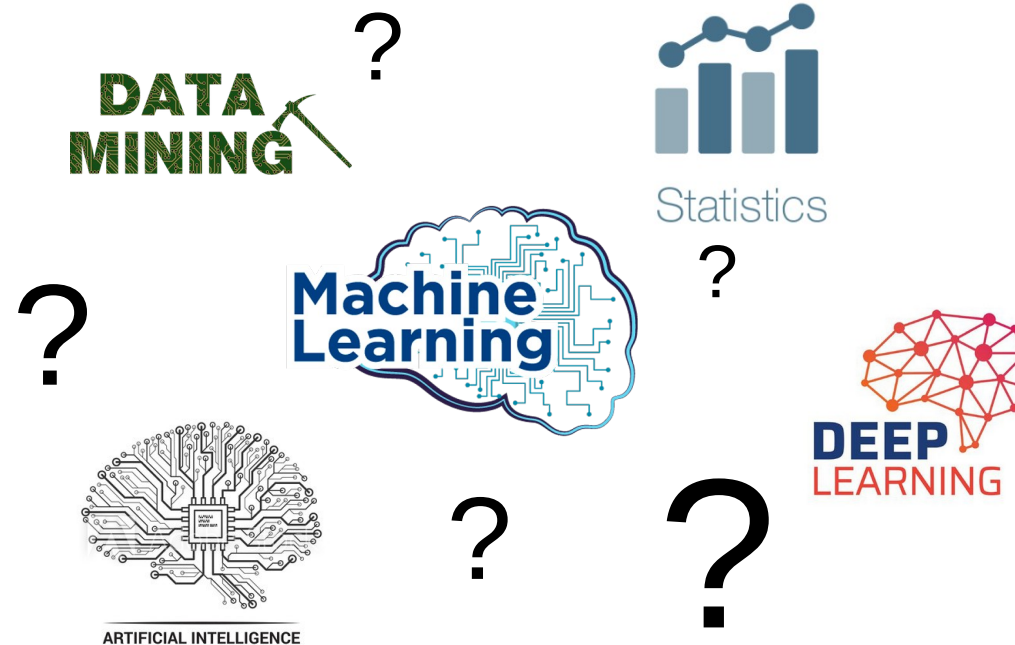
13:00 - 14:30 Ensemble methods (+ Interpretability)

14:30 - 15:15 Ensemble methods Exercise

15:15 - 16:15 Neural Networks

16:15 - 17:15 Neural Networks Exercise

What is Machine Learning?




- ▶ Distinction from Machine Learning to other statistical methodology not always clear
- ▶ When comparing Machine Learning with “classical” statistics:
 - ▶ Statistical models are generally designed for **inference**
 - ▶ Machine Learning models are generally designed for **prediction**

Application of Machine Learning

Being able to **predict** certain outcomes based on data can be important in many different areas in **research and industry**

Examples:



- ▶ Predict the winner of a basketball game
- ▶ Predict the weather of tomorrow
- ▶ Predict whether a medical scan shows an image of a tumor
- ▶ Predict whether an email is spam or not
- ▶ Predict how likely a person is about to develop depression



In all cases: **Predictions are based on data !**

Prediction models don't have to be complicated

- ▶ Simple linear regression can also be used to predict values of new observations




- ▶ However, sometimes statistical models have limited prediction accuracy, but allow **inference about the relation** between predictors and target variables (e.g. showing a significant influence of a treatment).
- ▶ In many Machine Learning models, the prediction accuracy is very good but it is difficult to interpret the variables' relations (e.g. neural network)

Application of Machine Learning

- ▶ Again: In general one tries to predict a target variable based on predictor variables


target variable ~ predictor variables
 $y \sim X$




- ▶ Target variable is usually a category or a number
 - ▶ Y is category: “Classification”
 - ▶ Y is metric: “Regression”
- ▶ In real-life data, there are often many predictor variables (genetic data: up to 10'000 predictors)
- ▶ Can even be $n \ll p$ (much more variables (p) than data points (n))
- ▶ This case can be difficult to handle with conventional methods (for example linear regression)

Challenges of high-dimensional data

- ▶ For example linear regression only works for $n > p$:



$n > p$




$n = 1$


- ▶ We need methods for situations with $n < p$
- ▶ Machine Learning methods are usually able to handle $n < p$ situations

Challenges of high-dimensional data


- ▶ For example linear regression only works for $n > p$:



$n > p$




$n = 2$




- ▶ We need methods for situations with $n < p$
- ▶ Machine Learning methods are usually able to handle $n < p$ situations


Outlook: Machine Learning methods




K-nearest neighbor




Principal Component Analysis




K-means clustering



Decision trees



Random Forest



Neural Networks

Quick R questionnaire

To get an impression of how used you are to R, think about the following statements:

- ▶ “I have never used R or R-Studio.”
- ▶ “I learned the basics of R once, but most of it is not really present anymore.”
- ▶ “I know the R basics.”
- ▶ “I am used to handling data sets in R and writing R scripts.”
- ▶ “I am used to writing `if()` statements and `for()` loops in R”
- ▶ “I use R every day.”

R – statistical computing

R is:

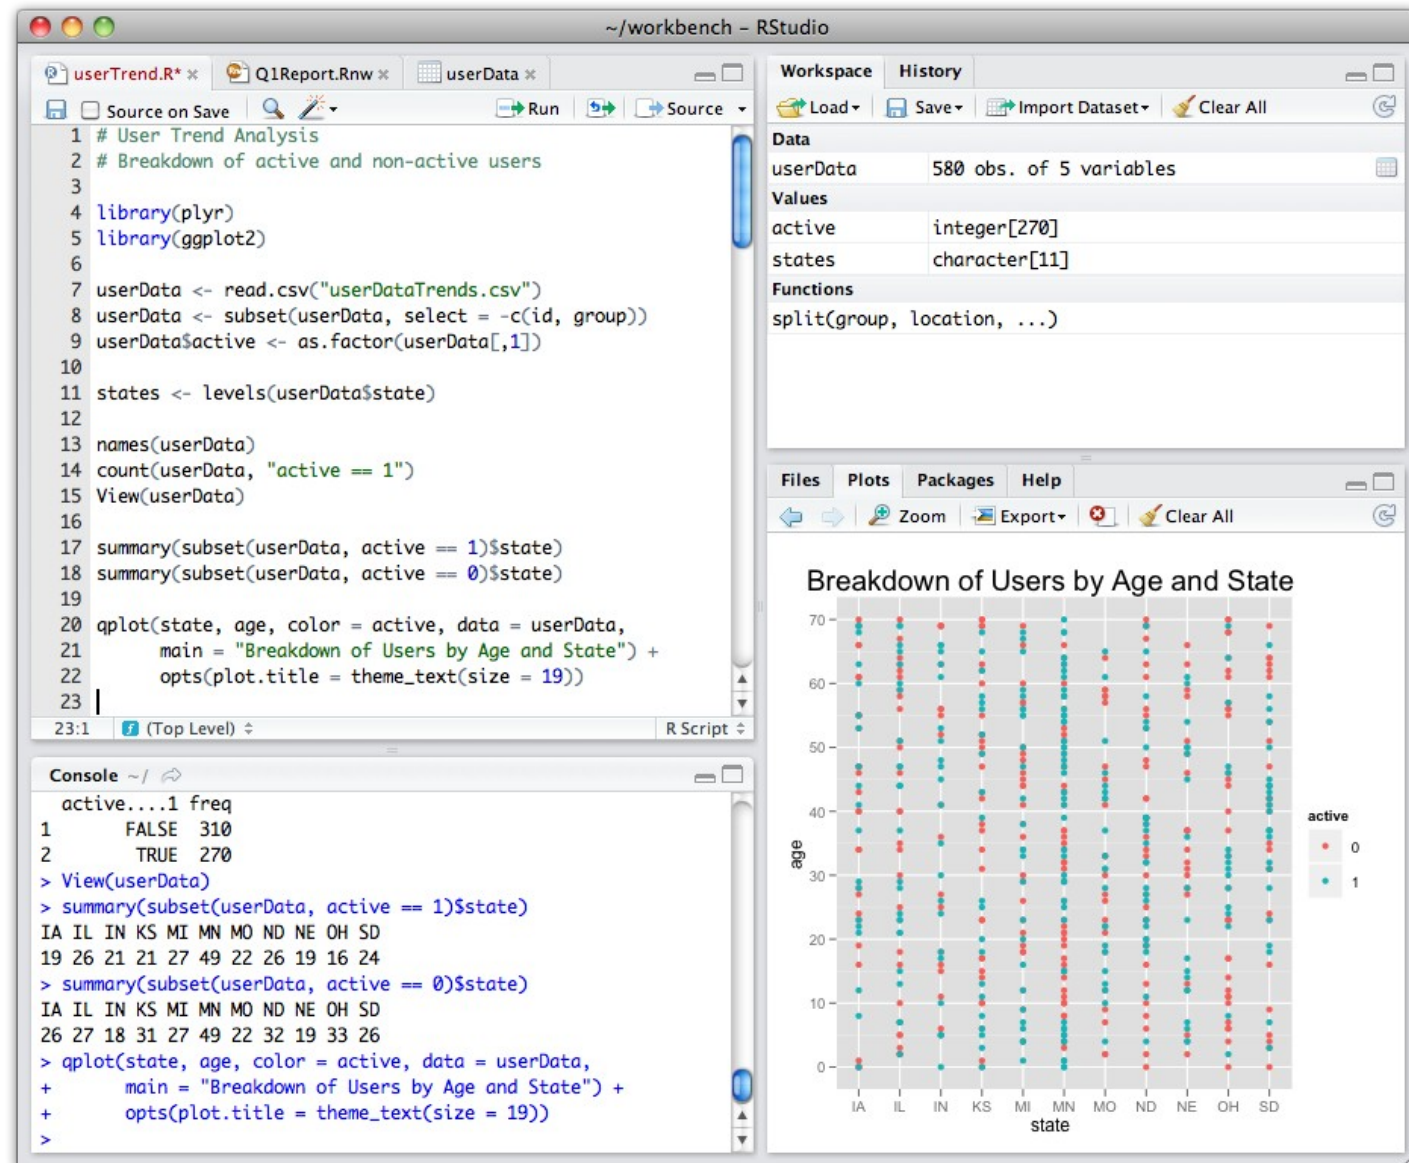
- ▶ Free to use, open-source
- ▶ Very flexible (> 10'000 add-on packages for R)
- ▶ Widely used among statisticians



In this lecture:

- ▶ Look at **R basics** (basic data types, simple functions, plotting, importing data,...)

R-studio: Integrated Development Environment (IDE) for R



R Introduction

- ▶ Classically we work in Rstudio in an R-script
- ▶ Execute the code from the script in the R-console
 - “source” to execute the whole script
 - Ctrl+enter to execute the current line or selection in script
- ▶ Using R as a calculator and to create objects:

```
> 1 + 1
```

```
[1] 2
```

```
> 10 * 30
```

```
[1] 300
```

```
> a <- 3 # create the variable a, which holds the value of 3
```

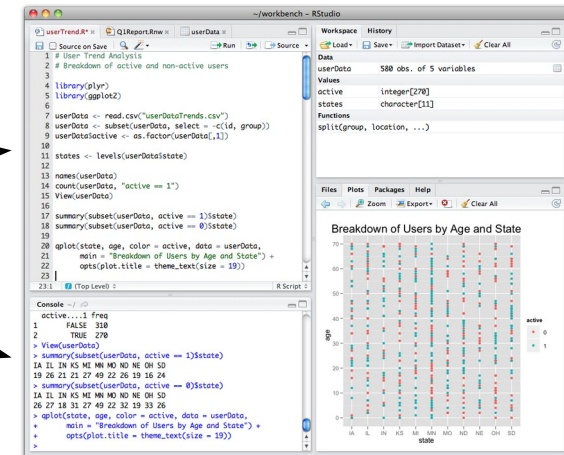
```
> b <- 25
```

```
> a + b
```

```
[1] 28
```

```
> a * b
```

```
[1] 75
```



Vectors (numerical and character)

- ▶ A vector is a combination of multiple elements and is created with `c()`:

```
> vec1 <- c(20, 122, 39)
```

```
> vec1
```

```
[1] 20 122 39
```

- ▶ A vector containing the integers from x to y can be created with `x:y`:

```
> 1:10
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

- ▶ Not only numerical values possible, example of a character-vector (text):

```
> vec2 <- c("Category1", "blue", "mixed")
```

```
> vec2
```

```
[1] "Category1" "blue" "mixed"
```

Selection of elements in vector

- ▶ When processing data one is often interested in selecting specific elements of the data. The selection of elements in R is performed with the square brackets `[]`:

```
# Selection of elements in vector:
```

```
> vecA <- c(2, 6, 7, 9)
```

```
> vecA
```

```
[1] 2 6 7 9
```

```
> vecA[2]
```

```
[1] 6
```

- ▶ Multiple elements can be selected by passing the wanted positions as a vector:

```
> vecA[c(2, 4)]
```

```
[1] 6 9
```

- ▶ Nested call:

```
> s <- c(2, 4)
```

```
> vecA[s]
```

```
[1] 6 9
```

- ▶ We can use negative indices to remove cases:

```
> vecA[-c(2, 4)]
```

```
[1] 2 7
```

```
> vecA[-s]
```

```
[1] 2 7
```

Data Frames (and reading in data)

- ▶ Data frames are commonly used to represent tabular data in R. The columns of a data frame are vectors. Exemplary data frame:

```
> dat
  vecA  vecB vecC
1     2   good 0.40
2     6   bad  0.20
3     7 medium 0.42
4     9   good 0.90
```

- ▶ Usually, a data frame is obtained by reading in a data file. In this workshop we will work with **.rda** files, .rda files can be read in with the **load()** command:

```
> load("toyDataFrame.rda")
> ls() # List objects in (global) environment to see what was loaded
```

- ▶ When reading in files as shown above, the **working directory** has to be set to the location of the file
 - ▶ The working directory is the folder where R looks for files to read in or where R saves created files (if not told otherwise)
 - ▶ The working directory can be set with **setwd(PATH TO DIRECTORY)** and viewed with **getwd()**. The working directory can also be set with the RStudio GUI (**Session > Set Working Directory**)

Selection of elements in a data frame

- ▶ In a data frame, columns can be selected using the \$ sign:

```
> dat
  vecA  vecB vecC
1    2   good 0.40
2    6   bad  0.20
3    7 medium 0.42
4    9   good 0.90

> dat$vecC
[1] 0.40 0.20 0.42 0.90

> dat$vecB
[1] "good" "bad" "medium" "good"

> dat$vecB[2]
[1] "bad"
```

- ▶ We can also use the \$ sign to add a new column or remove a column:

```
# Add new column:
> dat$newCol <- 1:2
> dat
  vecA  vecB vecC newCol
1    2   good 0.40      1
2    6   bad  0.20      2
3    7 medium 0.42      1
4    9   good 0.90      2
```

```
# Remove column:
> dat$vecB <- NULL
> dat
  vecA vecC newCol
1    2 0.40      1
2    6 0.20      2
3    7 0.42      1
4    9 0.90      2
```

Selection of elements in a data frame

- ▶ The square brackets can also be used to select elements in a data frame. In that case, the wanted row and column positions are passed to the brackets, separated by a comma:

```
> dat
```

```
  vecA  vecB vecC  
1    2   good 0.40  
2    6   bad  0.20  
3    7 medium 0.42  
4    9   good 0.90
```

```
# Element of dat in the second row  
in the third column:
```

```
> dat[2, 3]
```

```
[1] 0.20
```

```
# All elements in the first column:
```

```
> dat[, 1]
```

```
[1] 2 6 7 9
```

```
> dat[, "vecA"] # Selection with name
```

```
[1] 2 6 7 9
```

```
# First and fourth row with  
second and third column:
```

```
> dat[c(1, 4), c(2, 3)]
```

```
  vecB vecC  
1 good  0.4  
4 good  0.9
```

Logicals

- ▶ **Logicals** are an important class in R
- ▶ Logicals can only take the values **TRUE** or **FALSE**:

```
> LVec <- c(TRUE, TRUE, FALSE, TRUE)
> LVec
[1] TRUE TRUE FALSE TRUE
```

- ▶ **Logical operators** compare values:

```
> 5 > 3
[1] TRUE
> a <- 33
> b <- 5
> a == b # 'is equal to'
[1] FALSE
> c(6, 7, 2, 4, 1) < 3
[1] FALSE FALSE TRUE FALSE TRUE
```

Logicals can be used for selection

Selection with logicals:

```
> v <- 1:4
```

```
> v
```

```
[1] 1 2 3 4
```

```
> v[ c(TRUE, TRUE, FALSE, FALSE) ]
```

```
[1] 1 2
```

```
> dat
```

	vecA	vecB	vecC
1	2	good	0.40
2	6	bad	0.20
3	7	medium	0.42
4	9	good	0.90

```
> dat[ c(1,2), ]
```

	vecA	vecB	vecC
1	2	good	0.4
2	6	bad	0.2

```
> dat[ c(TRUE, TRUE, FALSE, FALSE), ]
```

	vecA	vecB	vecC
1	2	good	0.4
2	6	bad	0.2

```
> dat$vecA
```

```
[1] 2 6 7 9
```

```
> dat$vecA < 7
```

```
[1] TRUE TRUE FALSE FALSE
```

```
> dat[ dat$vecA < 7, ]
```

	vecA	vecB	vecC
1	2	good	0.4
2	6	bad	0.2

same thing

Categorical variables in R

- ▶ Categorical variables should ideally not be coded as numerical vectors in R
- ▶ There is the datatype **factor** specifically for categorical variables
- ▶ A factor not only contains the values of the individual elements, but also the list of possible categories ("levels")
- ▶ A factor can be created with **factor()**:

```
> d <- c(1, 2, 2, 3, 1)
> dfc <- factor(d, levels = c(1, 2, 3), labels = c("catA", "catB", "catC"))
> dfc
[1] catA catB catB catC catA
Levels: catA catB catC
```

Functions in R

- ▶ A function in R takes an **input**, processes it and returns an **output**
- ▶ Functions are applied by writing their name followed by normal brackets. The input for the function is defined in the brackets
- ▶ There are many ready-to-use functions in R:

Calculate the mean value of a vector:

```
> d <- c(1, 2, 2, 3, 1)
```

```
> mean(d)
```

```
[1] 1.8
```

Return the absolute value of a number:

```
> abs(-2)
```

```
[1] 2
```

Calculate the correlation coefficient of two vectors:


```
> cor(x = c(2,3,4,5,6), y = c(6,4,3,6,19), method = "spearman")
```

```
[1] 0.4616903
```

Take a random sample from a vector:

```
> sample(1:100, size = 5)
```

```
[1] 5 34 49 96 80
```



Functions in R

- ▶ How a function is applied (e.g. which input arguments exist, how they are named, ...) can be viewed by calling its help page (or searching the internet):

> ?mean

> ?cor

> ?sample

R Packages

- ▶ Publicly available R packages contain functions and objects for specific purposes
- ▶ R packages can be installed with **install.packages**("NAME OF PACKAGE") or via the RStudio GUI
- ▶ To make the content of a package readily available in an R session, the installed package has to be loaded with **library**(NAME OF PACKAGE)

Installing a package only has to be done once:

```
> install.packages("lme4") # Package for mixed models
```

Loading a package has to be done in each session:

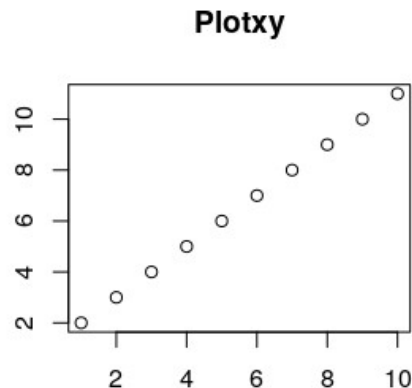
```
> library(lme4)
```

```
> lmer(y ~ ., data=d) # lmer() is a function from lme4
```

Creating plots in R


- ▶ There are many different ways and packages to create plots in R
- ▶ Basic scatterplots can be created with **plot()**:

```
> plot(x = 1:10, y = 2:11, main = "Plotxy")
```



- ▶ Boxplots can be created with **boxplot()**:

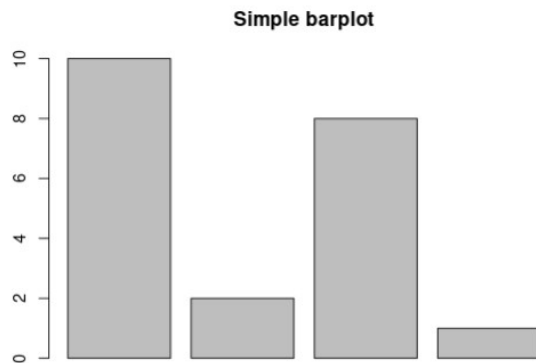
```
> boxplot(rnorm(30), main='boxplot') # rnorm(30) draws 30 datapoints from  
                                     # the (standard) normal distribution
```



Creating plots in R

- ▶ There are many different ways and packages to create plots in R
 - ▶ Barplot created with **barplot()**:

```
> barplot( c(10, 2, 8, 1) , main = "Simple barplot")
```



Two basic programming structures

- ▶ If()-statements
- ▶ for()-loops

If() statements

- ▶ If-statements are a very important tool in programming
- ▶ An if-statement tells R to execute a block of commands, if a condition is true
- ▶ Structure is always the same:
 - ▶ **if** (**condition**) {*execute if TRUE*} **else** {*execute if FALSE*}
 - ▶ Simple example:

```
a <- 2
b <- 5
if (a < b){
  a <- a + 1 # execute if condition is true
}else{
  a <- a - 1 # execute if condition is false
}

> a
[1] 3 # 2 + 1
```

logical condition
(TRUE in this case)

```
# reversed case:
a <- 20
b <- 5
if (a < b){
  a <- a + 1
}else{
  a <- a - 1 # gets executed
}

> a
[1] 19 # 20 - 1
```

FALSE

If() statements

- ▶ One can also use more complicated conditions:
 - ▶ If($a > b$ & $a > c$){...} “execute code if a is bigger than b **and** a is bigger than c”
 - ▶ If($a > b$ | $a > c$){...} “execute code if a is bigger than b **or** a is bigger than c”
- ▶ **Example:** Use if-statement to see if there are **invalid values** in a vector
- ▶ For a collected variable it is known that values larger than 10 are not possible (e.g. questionnaire where 10 is maximum answer)

```
> questionnaire_data
[1] 6 11 3 1 6 5 3 6 1 5 3 3 9 5 7 5 3 3 3 4
> if(max(questionnaire_data) > 10){
  cat('there are invalid entries in this data') # execute if true
} else {
  cat('no invalid data') # execute if false
}
there are invalid entries in this data
```

If() statements

- ▶ One can also combine multiple if statements using *else if*:

```
if (condition1) {  
    execute this  
} else if (condition2) {  
    execute this  
} else {  
    execute this  
}
```

For() loops

- ▶ **For**-loops are another important tool in programming
- ▶ For-loops execute a selection of code multiple times
- ▶ The structure in R looks as follows:
 - ▶ for(**variable** in **vector**) { *execute code for variable* }
 - ▶ Simple example:

```
> for(i in c(2,5,6,7)){ # i is variable, c(2,5,6,7) is the vector
  a <- i + 5
  print(a) # force to show a
}
```

[1] 7
[1] 10
[1] 11
[1] 12

For() loops

- ▶ **Good trick:** use the variable of the for-loop as an **index**:

```
> a <- NA # declare "a" because we use a[] later
> for(i in 1:5){
  a[i] <- i + 5
}
> a
[1] 6 7 8 9 10
```

- ▶ **If-statements** and **for-loops** are often combined with each other

- ▶ E.g. Double only the values smaller than 20 in a vector:

```
> a <- c(2, 44, 6, 4.5, 94)
> a
[1] 2.0 44.0 6.0 4.5 94.0
> b <- NA
> for (i in 1:length(a)){ # same as 1:5
  if(a[i] < 20){b[i] <- a[i]*2} else {
    b[i] <- a[i]
  }
}
```

```
> b
[1] 4 44 12 9 94
```